

ML Lab 3 Answers

1. Reducing to 100 Iterations

Observation:

- Initial cost (at $\theta=[0;0]$) remains ~ 32.07
- Final θ values show minimal change from initial $[0;0]$
- Linear regression line stays nearly flat
- Predictions become less accurate:

For population=35,000: $\sim \$0$ -5,000 profit

For population=70,000: $\sim \$0$ -10,000 profit ``

Explanation:

Gradient descent hasn't had enough steps to:

- Properly descend the cost function surface
- Adjust θ parameters meaningfully
- Reach the convex bowl of the cost function

2. Increasing to 10,000 Iterations

Observation:

- Cost decreases further (beyond original 1500-iteration value)
- θ converges closer to optimal values:

$\theta \approx [-3.63; 1.16]$ (vs original $[-3.88; 1.19]$) ``

- Regression line fits data better
- Predictions become more stable:

For 35,000: $\sim \$45,000$ profit

For 70,000: $\sim \$91,000$ profit ``

Explanation:

Additional iterations allow:

- More precise convergence to minimum
- Better parameter optimization
- Smoother traversal down cost surface
- Stabilization of gradient updates

Key Visual Differences

1. **Cost Function Plot:**

- 100 iterations: Contour plot shows theta far from center
- 10k iterations: Theta marker reaches deepest contour level

2. **Learning Curve:**

(Plot J_history vs iterations)

- 100 iterations: Curve still descending sharply
- 10k iterations: Curve flattens showing convergence

Technical Insight

The gradient descent update rule in gradientDescent.m:

$\theta(1) = \theta(1) - (\alpha/m) * \text{temp0};$

$\theta(2) = \theta(2) - (\alpha/m) * \text{temp1};$

Works better with:

- **More iterations:** Allows smaller, precise steps near minimum
- **Learning rate ($\alpha=0.01$):** Prevents overshooting with high iterations

Pro Tip: For this dataset, 1500-3000 iterations with $\alpha=0.01$ typically suffices. Extreme values (>1M iterations) show diminishing returns due to floating-point precision limits.